

Practical-5 CBOW

Imports and Initial Setup

```
```python
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib as mpl
import matplotlib.pylab as pylab
import numpy as np
%matplotlib inline
```
```

Explanation:

- **matplotlib** and **seaborn** are for visualization. **matplotlib.pyplot** provides plotting functions, and **seaborn** offers advanced styling for plots.
- **%matplotlib inline** enables inline plotting in Jupyter notebooks (if you're using one).
- **numpy (np)** is for numerical operations, especially useful for handling arrays and matrices, which are crucial in machine learning.

Data Preparation

```
```python
import re
```
```

Explanation:

- **re**: This library handles regular expressions, which are useful for text preprocessing, like removing unwanted characters.

```
```python
sentences = """We are about to study the idea of a computational process.
Computational processes are abstract beings that inhabit computers.
As they evolve, processes manipulate other abstract things called data.
The evolution of a process is directed by a pattern of rules
called a program. People create programs to direct processes. In effect,
we conjure the spirits of the computer with our spells."""
```
```

Explanation:

- This variable **sentences** contains a block of text used for training the CBOW model. It's meant to simulate a small corpus or set of sentences, which the model will learn word relationships from.

```
```python
```

```
Clean Data
```

```
remove special characters
```

```
sentences = re.sub('[^A-Za-z0-9]+', ' ', sentences)
```

```
remove 1 letter words
```

```
sentences = re.sub(r'(?!\w)(?!\s)', ' ', sentences).strip()
```

```
lower all characters
```

```
sentences = sentences.lower()
```

```
```
```

Explanation:

- **Data Cleaning**:

- **Remove special characters**: Replaces any non-alphanumeric characters with spaces to clean up the text.

- **Remove single characters**: Gets rid of isolated single characters, as they may not carry useful information in this context.

- **Convert to lowercase**: Makes the text lowercase to ensure that words like "Process" and "process" are treated the same.

Vocabulary Preparation

```
```python
```

```
words = sentences.split()
```

```
vocab = set(words)
```

```
```
```

Explanation:

- **Tokenization**:

- **words**: Splits the cleaned text into a list of individual words.

- **vocab**: Creates a set of unique words, which will be used to build the vocabulary. This set will be used to assign each word a unique index for training.

```
```python
```

```
vocab_size = len(vocab)
embed_dim = 10
context_size = 2
```
```

Explanation:

- **Define Model Parameters**:
 - **vocab_size**: Total number of unique words in our vocabulary, needed for building word embeddings.
 - **embed_dim**: Dimensionality of the word embeddings. Each word will be represented as a vector of 10 values.
 - **context_size**: Number of words to look at on each side of a target word. Here, we look 2 words before and 2 words after.

Vocabulary Indexing

```
```python
word_to_ix = {word: i for i, word in enumerate(vocab)}
ix_to_word = {i: word for i, word in enumerate(vocab)}
```
```

Explanation:

- **Create Word-Index Mappings**:
 - **word_to_ix**: Maps each word in the vocabulary to a unique integer index.
 - **ix_to_word**: Reverse mapping to retrieve the word given its index. This is useful for converting predictions back into words.

Data Preparation for CBOW

```
```python
Data bags
data - [(context), target]
data = []
for i in range(2, len(words) - 2):
 context = [words[i - 2], words[i - 1], words[i + 1], words[i + 2]]
 target = words[i]
 data.append((context, target))
print(data[:5])
```
```

Explanation:

- **Generate Training Pairs**:
 - **Loop through each word** with a range starting from index 2 and ending

two indices before the last word.

- **Context words**: For each target word, define its context by selecting two words before and two words after it.
- **Target word**: The word we want to predict based on its context.
- **Append (context, target) pairs** to `data`. Each pair will be used for training the CBOW model.
- **print(data[:5])**: Prints the first five pairs for verification.

Initialize Embeddings

```
```python
embeddings = np.random.random_sample((vocab_size, embed_dim))
```
```

Explanation:

- **Initialize Random Embeddings**:
 - **embeddings**: Creates an array of random numbers with dimensions `vocab_size x embed_dim`. Each row represents a word's embedding vector.
 - This is the initial random embedding matrix, which will be adjusted during training to capture word meanings.

Linear Model

```
```python
Linear Model
def linear(m, theta):
 w = theta
 return m.dot(w)
```
```

Explanation:

- **Define Linear Transformation**:
 - **linear function**: Takes an input matrix `m` and applies a transformation using the weights matrix `theta` (also called `w` here).
 - **dot product**: Multiplies input matrix `m` with weights `w` to obtain transformed output, which will later be passed to an activation function.

Log Softmax and Cross Entropy Loss

```
```python
Log softmax + NLLloss = Cross Entropy
```

```
def log_softmax(x):
 e_x = np.exp(x - np.max(x))
 return np.log(e_x / e_x.sum())
````
```

Explanation:

- **Log Softmax**:
- **log_softmax** function computes the log of softmax values for an input `x`.
- **Softmax**: Transforms the output values into probabilities, with each value representing the likelihood of a word being the correct target.
- **Log**: Converts probabilities into log-probabilities, which are easier to use with the cross-entropy loss.

```
````python
def NLLLoss(logs, targets):
 out = logs[range(len(targets)), targets]
 return -out.sum()/len(out)
````
```

Explanation:

- **Negative Log-Likelihood Loss (NLLLoss)**:
- **logs**: Log-probabilities from `log_softmax`.
- **targets**: Indices of actual target words.
- **Select target log-probabilities**: Extracts the log-probabilities of correct target words.
- **Mean Negative Log-Sum**: Takes the average of the negative log-sum, giving the model a measure of its error. This is what the model tries to minimize during training.

```
````python
def log_softmax_crossentropy_with_logits(logits, target):
 out = np.zeros_like(logits)
 out[np.arange(len(logits)), target] = 1
 softmax = np.exp(logits) / np.exp(logits).sum(axis=-1, keepdims=True)
 return (- out + softmax) / logits.shape[0]
````
```

Explanation:

- **Combined Log Softmax and Cross Entropy**:
- **One-hot encoding** for `target`: Converts the target indices into a format where only the correct class is 1, and others are 0.
- **Softmax** computes probabilities, similar to `log_softmax`, but without

taking the log.

- **Error Computation**: Returns the gradient of the cross-entropy loss with respect to logits for backpropagation.

Forward Pass Function

```
```python
def forward(context_idxxs, theta):
 m = embeddings[context_idxxs].reshape(1, -1)
 n = linear(m, theta)
 o = log_softmax(n)
 return m, n, o
```
```

Explanation:

- **Forward Pass**:
 - **context_idxxs**: The indices of context words.
 - **m**: Retrieves and flattens embeddings for the context words.
 - **n**: Applies the linear transformation to get logits.
 - **o**: Applies log_softmax to get log-probabilities, which are used to calculate the loss.

Backward Pass (Gradient Calculation)

```
```python
def backward(preds, theta, target_idxxs):
 m, n, o = preds
 dlog = log_softmax_crossentropy_with_logits(n, target_idxxs)
 dw = m.T.dot(dlog)
 return dw
```
```

Explanation:

- **Backward Pass**:
 - **preds**: Output from the forward pass.
 - **dlog**: Computes gradients with respect to the logits.
 - **dw**: Calculates gradient of the weights `theta` by multiplying context embedding matrix `m.T` with `dlog`.

Optimization Step

```

```python
def optimize(theta, grad, lr=0.03):
 theta -= grad * lr
 return theta
```

```

Explanation:

- **Update Weights**:
- **theta**: Weight matrix for the model.
- **grad**: Gradient of the loss with respect to `theta`.
- **lr**: Learning rate, controlling the update size.
- **theta -= grad * lr**

Training Loop: Generate Training Data and Optimize Weights

```

```python
Training
#Generate training data
theta = np.random.uniform(-1, 1, (2 * context_size * embed_dim, vocab_size))
```

```

Explanation:

- **Initialize the `theta` matrix**:
- **theta**: A randomly initialized weight matrix that will be updated during training.
- **Dimensions**: `(2 * context_size * embed_dim, vocab_size)` means this matrix has rows equal to the number of input features (determined by `context_size` and `embed_dim`) and columns equal to the vocabulary size. This structure helps the model map from input embeddings to output probabilities.

```

```python
epoch_losses = {}

for epoch in range(80):

 losses = []

 for context, target in data:
 context_idx = np.array([word_to_ix[w] for w in context])
 preds = forward(context_idx, theta)

```

```

 target_idx = np.array([word_to_ix[target]])
 loss = NLLLoss(preds[-1], target_idx)

 losses.append(loss)

 grad = backward(preds, theta, target_idx)
 theta = optimize(theta, grad, lr=0.03)

 epoch_losses[epoch] = losses
 ...

```

#### #### Explanation:

- **Training Loop**:
  - **epoch\_losses**: Dictionary to store the loss for each epoch.
  - **epochs (80)**: The model goes through the entire dataset 80 times to learn patterns in the data.
- **Inner Loop (per context-target pair)**:
  - **context\_idx**: Converts context words into their respective indices.
  - **forward**: Computes the predicted output (``preds``) for the context words using the ``theta`` matrix.
  - **target\_idx**: Converts the target word into its index.
  - **loss**: Calculates the error between the predicted output and the actual target using Negative Log-Likelihood Loss (``NLLLoss``).
  - **backward**: Computes gradients of the loss with respect to ``theta``.
  - **optimize**: Updates ``theta`` by applying gradients, allowing the model to improve its predictions.

---

#### ### Loss Analysis: Plotting Loss per Epoch

```

```python
# Analyze
# Plot loss/epoch
ix = np.arange(0,80)
fig = plt.figure()
fig.suptitle('Epoch/Losses', fontsize=20)
plt.plot(ix,[epoch_losses[i][0] for i in ix])
plt.xlabel('Epochs', fontsize=12)
plt.ylabel('Losses', fontsize=12)
```

```

#### #### Explanation:

- **Loss Plotting**:
  - **ix**: Array of epoch numbers from 0 to 80.



- **plt.plot**: Plots the loss per epoch. This shows how the model's error decreases as it trains, indicating it is learning.
- **xlabel and ylabel**: Labeling the axes to show "Epochs" on the x-axis and "Losses" on the y-axis for easy interpretation.

---

### ### Prediction Function

```
```python
# Predict function
def predict(words):
    context_idx = np.array([word_to_ix[w] for w in words])
    preds = forward(context_idx, theta)
    word = ix_to_word[np.argmax(preds[-1])]
    return word
```
```

#### #### Explanation:

- **Predict Function**:
  - **context\_idx**: Converts the input words (context) into indices using `word_to_ix`.
  - **forward**: Generates predictions using the current `theta` matrix.
  - **np.argmax**: Selects the word with the highest probability as the predicted target word.
  - **ix\_to\_word**: Converts the index back to the actual word using `ix_to_word`.
- This function allows us to test the model by inputting a few context words and seeing what target word it predicts.

---

### ### Example Prediction

```
```python
# (['we', 'are', 'to', 'study'], 'about')
predict(['we', 'are', 'to', 'study'])
```
```

#### #### Explanation:

- **Example Use of `predict`**:
  - This line predicts the target word for the context words `['we', 'are', 'to', 'study']`.
  - **Expected Output**: The model will output the word it finds most probable to be associated with this context, helping us check if the model learned meaningful relationships.

---

### ### Model Accuracy Function

```
```python
# Accuracy
def accuracy():
    wrong = 0

    for context, target in data:
        if(predict(context) != target):
            wrong += 1

    return (1 - (wrong / len(data)))
```
```

#### #### Explanation:

- **Accuracy Calculation**:
  - Loop through all context-target pairs in `data`.
  - **Prediction Check**: If the predicted target word doesn't match the actual target, increment the `wrong` counter.
  - **Accuracy Formula**: `(1 - (wrong / len(data)))` gives the percentage of correct predictions. Higher accuracy means the model has learned well.

---

### ### Final Prediction Example

```
```python
predict(['processes', 'manipulate', 'things', 'study'])
```
```

#### #### Explanation:

- **Another Example Prediction**:
  - This line tests the model by inputting the context `['processes', 'manipulate', 'things', 'study']`.
  - It checks whether the model can predict a meaningful target word based on this new context, allowing us to evaluate the learned associations.

---

Here are key questions for Assignment 5, which focuses on implementing the Continuous Bag of Words (CBOW) model for natural language processing (NLP):

### ### Important Questions and Answers

1. **What is Natural Language Processing (NLP)?**
  - NLP is a field of AI that enables computers to understand, interpret, and generate human language.
2. **What is Word Embedding in NLP?**
  - Word embeddings are vector representations of words that capture semantic relationships, allowing similar words to have similar vector representations.
3. **What is the Word2Vec technique?**
  - Word2Vec is a model that transforms words into vector space, using CBOW or Skip-gram to learn word relationships based on context.
4. **Explain the architecture of the Continuous Bag of Words (CBOW) model.**
  - In CBOW, the model predicts a target word from a given context of surrounding words, learning embeddings for each word based on this context.
5. **What is the input and output of the CBOW model?**
  - **Input**: Context words (surrounding words in a sentence).
  - **Output**: Target word (the word predicted based on the context).
6. **What is the purpose of a Tokenizer in NLP?**
  - A tokenizer splits text into smaller units (tokens), like words or sentences, enabling easier processing for NLP tasks.
7. **Explain the window size parameter in the CBOW model.**
  - Window size determines the number of context words used on each side of the target word, affecting how much surrounding context the model considers.
8. **What are Embedding and Lambda layers in Keras?**
  - **Embedding Layer**: Transforms integer-encoded words into dense vectors.
  - **Lambda Layer**: Allows custom operations to be applied within the model, often used for function-based transformations.
9. **What is the purpose of the yield() function in Python?**
  - `yield()` is used in generators to return intermediate values without exiting the function, allowing iteration over large data without excessive memory use.
10. **Describe the process of data preparation for training a CBOW model.**
  - Data preparation involves tokenizing text, creating word pairs of context and target words, and encoding them into numerical format for the model.
11. **Why is the Sequential model commonly used in CBOW implementations?**

- Sequential models are simple to set up and suitable for layer-wise stacking, making them ideal for straightforward architectures like CBOW.

12. **What is Gensim, and how is it used in NLP?**

- Gensim is an NLP library in Python for topic modeling and word embedding. It provides tools to train models like Word2Vec, CBOW, and Skip-gram.

13. **How does the CBOW model represent word similarity?**

- The CBOW model learns word embeddings such that similar words have similar vector representations, allowing the model to predict related words accurately.

14. **Explain the concept of context and target words in CBOW.**

- **Context Words**: Words surrounding a target word in a sentence.
- **Target Word**: The word being predicted based on its surrounding context.

15. **What is the purpose of vectorization in CBOW?**

- Vectorization converts words into numerical representations (embeddings), enabling the model to perform mathematical operations and learn relationships.

Would you like additional questions or more details on any of these answers?