

Practical-6 object detection

Importing Libraries

```
```python
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import numpy as np
from tensorflow.keras.applications import VGG16
```
```

****Explanation**:**

- ****TensorFlow and Keras****: Used to create and train the neural network.
- ****VGG16****: Pre-trained VGG16 model, which we'll use as a base for our custom classifier.
- ****Model, Dense, Flatten****: These are Keras layers and model structures for creating and modifying neural networks.
- ****ImageDataGenerator****: Used to preprocess and augment image data for training and testing.
- ****NumPy****: A library for handling arrays and numerical data.

Define Training and Testing Data Directories

```
```python
train_dir = r"C:\Users\Atharva\Desktop\apudl\1st\cifar-10-img\cifar-10-
img\train"
test_dir = r"C:\Users\Atharva\Desktop\apudl\1st\cifar-10-img\cifar-10-
img\test"
```
```

****Explanation**:**

- ****train_dir**** and ****test_dir****: These are paths to the training and testing datasets. Each directory contains images for CIFAR-10 classes.

Data Preprocessing with ImageDataGenerator

```

```python
train_datagen = ImageDataGenerator(
 rescale=1.0 / 255,
)
test_datagen = ImageDataGenerator(
 rescale=1.0 / 255,
)
```

```

****Explanation**:**

- ****ImageDataGenerator****: This tool helps to prepare image data before feeding it to the model.
- ****rescale****: Divides each pixel value by 255 to normalize them to the range [0, 1]. This helps the model learn more efficiently.

Create Training and Testing Generators

```

```python
train_batch_size = 5000
train_generator = train_datagen.flow_from_directory(
 train_dir,
 target_size=(32, 32),
 batch_size=train_batch_size,
 class_mode='categorical'
)
test_batch_size = 1000
test_generator = test_datagen.flow_from_directory(
 test_dir,
 target_size=(32, 32),
 batch_size=test_batch_size,
 class_mode='categorical'
)
```

```

****Explanation**:**

- ****train_generator**** and ****test_generator****: These generators load images from the directories and preprocess them.
- ****target_size****: Resizes all images to 32x32 pixels (CIFAR-10 standard).
- ****batch_size****: Number of images loaded per batch.
- ****class_mode='categorical'****: Specifies that classes are categorical (10 classes in CIFAR-10).

Load Sample Data from Generators

```
```python
x_train, y_train = train_generator[0]
x_test, y_test = test_generator[0]
print(len(x_train))
print(len(x_test))
```
```

****Explanation**:**

- ****x_train, y_train****: Load one batch of training images and labels from ``train_generator``.
- ****x_test, y_test****: Load one batch of testing images and labels from ``test_generator``.
- ****print(len(x_train))****: Prints the number of images in this batch to confirm the data loading.

Load Pre-Trained VGG16 Model and Freeze Layers

```
```python
weights_path = r"C:\Users\Atharva\Desktop\apudl\1st\cifar-10-img\cifar-10-
img\vgg16_weights_tf_dim_ordering_tf_kernels_notop (2).h5"
base_model = VGG16(weights=weights_path, include_top=False,
input_shape=(32, 32, 3))

for layer in base_model.layers:
 layer.trainable = False
```
```

****Explanation**:**

- ****base_model****: Loads the VGG16 model without the top classification layer (``include_top=False``), allowing us to customize the final layers for CIFAR-10 classification.
- ****input_shape=(32, 32, 3)****: Sets the input size to match CIFAR-10 images (32x32 pixels, 3 color channels).
- ****weights=weights_path****: Uses the pre-trained weights stored at the specified path.
- ****Freeze layers****: Sets ``layer.trainable = False`` for each layer, so the base layers won't be updated during training, saving time and retaining VGG16's learned features.

Add Custom Classification Layers

```

```python
x = Flatten()(base_model.output)
x = Dense(256, activation='relu')(x)
x = tf.keras.layers.Dropout(0.3)(x)
x = Dense(256, activation='relu')(x)
x = tf.keras.layers.Dropout(0.3)(x)
predictions = Dense(10, activation='softmax')(x)
model = Model(inputs=base_model.input, outputs=predictions)
model.compile(optimizer="adam", loss='categorical_crossentropy',
metrics=['accuracy'])
```

```

****Explanation**:**

- ****Flatten****: Converts the VGG16 output (a 3D feature map) to a 1D vector.
- ****Dense(256, activation='relu')****: Adds a fully connected layer with 256 units and ReLU activation for non-linear learning.
- ****Dropout(0.3)****: Randomly turns off 30% of the neurons to prevent overfitting.
- ****Dense(10, activation='softmax')****: Final layer with 10 units (one per class), using softmax activation to output probabilities for each class.
- ****model.compile****: Configures the model for training with:
 - ****optimizer="adam"****: Optimization algorithm.
 - ****loss='categorical_crossentropy'****: Suitable for multi-class classification.
 - ****metrics=['accuracy']****: Tracks model accuracy during training.

Train the Model

```

```python
model.fit(x_train, y_train, batch_size=64, epochs=10, validation_data=(x_test,
y_test))
```

```

****Explanation**:**

- ****model.fit****: Trains the model on training data.
 - ****batch_size=64****: Trains on 64 samples at a time.
 - ****epochs=10****: Iterates through the data 10 times.
 - ****validation_data=(x_test, y_test)****: Evaluates accuracy on test data after each epoch.

Make Predictions

```

```python

```

```
import matplotlib.pyplot as plt
predicted_value = model.predict(x_test)
````
```

****Explanation**:**

- ****predicted_value****: Uses the trained model to predict labels for `x_test`.
- ****plt****: Imports the `matplotlib.pyplot` library for displaying images and predictions later.

Load Class Labels

```
``python
labels = list(test_generator.class_indices.keys())
````
```

**\*\*Explanation\*\*:**

- **\*\*labels\*\***: Retrieves the list of class names from the test data generator. This will help display actual class names instead of just numerical labels.

---

### ### Display Predictions

```
``python
n=234
plt.imshow(x_test[n])
print("Predicted: ", labels[np.argmax(predicted_value[n])])
print("Actual: ", labels[np.argmax(y_test[n])])
````
```

****Explanation**:**

- ****plt.imshow(x_test[n])****: Displays the nth image from the test set.
- ****np.argmax(predicted_value[n])****: Finds the class with the highest probability for the nth image.
- ****np.argmax(y_test[n])****: Retrieves the actual class label for comparison.

The code repeats this for different values of `n` to display more predictions.

Here are the key questions for Assignment 6, which focuses on object detection using transfer learning with Convolutional Neural Network (CNN) architectures:

Important Questions and Answers

1. **What is Transfer Learning?**

- Transfer learning is a technique in deep learning where a model trained on one task is adapted for a similar but different task, saving time and resources.

2. **What are Pretrained Neural Network models?**

- Pretrained models are models that have already been trained on large datasets, such as ImageNet, and can be fine-tuned for specific tasks.

3. **Explain the advantages of Transfer Learning.**

- Transfer learning accelerates training, reduces the need for large datasets, and improves model accuracy by leveraging knowledge from similar tasks.

4. **List some applications of Transfer Learning.**

- Image classification, object detection, natural language processing, and medical image analysis.

5. **What is the Caltech 101 dataset?**

- Caltech 101 is an image dataset containing pictures of objects from 101 categories, often used for object recognition research.

6. **Explain the ImageNet dataset.**

- ImageNet is a large dataset with over 14 million images across 20,000 categories, commonly used for training deep learning models.

7. **List the basic steps in Transfer Learning.**

- Load a pretrained model, freeze the initial layers, add a custom classifier, fine-tune with new data, and evaluate performance.

8. **What is Data Augmentation, and why is it used in Transfer Learning?**

- Data augmentation involves creating modified versions of the training data (e.g., rotations, flips) to increase dataset diversity, improving model generalization.

9. **Why is preprocessing needed for input data in Transfer Learning?**

- Preprocessing standardizes inputs, ensuring compatibility with the pretrained model and improving training stability and accuracy.

10. **What is the PyTorch Transforms module, and what are some common transformations?**

- The PyTorch Transforms module provides functions to manipulate image data, such as resizing, cropping, flipping, and normalization, which prepare data for CNNs.

11. **Explain the purpose of the `Compose` function in PyTorch Transforms.**

- `Compose` chains multiple transformations sequentially, applying them to the input data in a defined order.

12. ****What is VGG-16, and why is it used in Transfer Learning?****

- VGG-16 is a CNN model with 16 layers, known for its simplicity and effectiveness in image classification. It is widely used for transfer learning due to its architecture.

13. ****What does it mean to "freeze" model parameters in Transfer Learning?****

- Freezing parameters prevents certain layers from updating during training, preserving the learned features from the pretrained model.

14. ****Why is a custom classifier added on top of a pretrained model in Transfer Learning?****

- A custom classifier tailors the model to the specific task, as the original model's classifier may not match the number of classes in the new task.

15. ****What is Early Stopping, and how is it used in Transfer Learning?****

- Early stopping halts training when performance stops improving on the validation set, preventing overfitting and saving resources.